

nufftcf: Fast Auto- and Cross-Correlation Function Estimation for Irregularly-Sampled Time Series via the Non-Uniform FFT

Jean-Eric Campagne 

Université Paris-Saclay, CNRS/IN2P3, IJCLab, 91405 Orsay, France
`jean-eric.campagne@ijclab.in2p3.fr`

September 4, 2026

Abstract

Estimating the auto-correlation function (ACF) and the cross-correlation function (CCF) of sampled series is a standard task across different physics fields as e.g. astronomy, and environmental fields. We aim to provide an ACF/CCF estimator primarily designed for irregularly sampled series that is numerically consistent with established, well-validated kernel-weighted definitions, while scaling as $O(n \log n)$ rather than $O(n^2)$.

nufftcf evaluates the Wiener–Khinchin theorem for irregularly sampled data using the Non-Uniform Fast Fourier Transform (NUFFT), via the Flatiron Institute’s **FINUFFT** library, and using the Gaussian- and rectangle-kernel estimators. An $O(n)$ two-pointer scan replaces the naive $O(n^2)$ computation of the effective pair count that normalizes each lag bin. Direct, real-space implementations of the same estimators for small series and for reference, and dedicated classical-FFT estimators for regularly-sampled data, are provided alongside the NUFFT implementation, all sharing a single calling convention.

With synthetic times series, we validate **nufftcf** for ACF against **pastas**, a library known in ground-water time series analysis, and for CCF against **pyZDCF**, known in astronomical time series analysis. Benchmarks confirm the expected asymptotic scalings – $O(n \log n)$ for **nufftcf** versus $O(n^2)$ for **pastas** slotting technique – with **nufftcf** already advantageous at moderate series lengths thanks to its low (millisecond-scale) overhead. On regularly-sampled data, **nufftcf** dedicated FFT path further reduces both cost and overhead relative to its own NUFFT estimator. Beyond this cross-validation exercise, we further demonstrate **nufftcf** on a simulated ground-based stellar light curve combining a quasi-periodic rotation signal, correlated flicker noise, seasonal sampling gaps, and heteroscedastic measurement errors.

Keywords: Time series analysis, Correlation function, Non-uniform Fast Fourier Transform

1 Introduction

Many measurements of physical, environmental, and astrophysical systems take the form of a scalar time series that is not regularly sampled in time. Ground-based photometric monitoring of active galactic nuclei (AGN) is interrupted by weather, daylight, and telescope scheduling; groundwater levels are logged during irregular field visits; paleoclimate proxies are

deposited at a rate that itself varies with the climate signal being measured. A recurring analysis task for such data is to estimate the autocorrelation function (ACF) of a single series, or the cross-correlation function (CCF) between two independently and irregularly sampled series, at a set of time lags τ .

Applications range from measuring a characteristic variability or persistence timescale, to detecting a time delay between two related signals. For instance

in AGN reverberation mapping, where the CCF between a driving continuum light curve and a reprocessed emission-line light curve encodes the size of the emitting region (W. Welsh, 2024; I. Jankov et al., 2022a; P. Sánchez-Sáez et al., 2021; A. I. Shapovalova et al., 2019; A. Kovačević et al., 2014), one uses the discrete correlation function (R. A. Edelson & J. H. Krolik, 1988), and also for instance the python version (`pyZDCF`) of the code by T. Alexander (1997). A comparable need arises in hydro(geo)logical time series analysis: the open-source Python package `Pastas` (R. A. Collenteur et al., 2019), widely used for time series modeling of groundwater levels, implements an ACF estimator for irregularly sampled series based on K. Rehfeld et al. (2011) and R. A. Edelson & J. H. Krolik (1988).

Because the classical, textbook estimator of the ACF/CCF assumes a common, regular sampling grid (e.g. D. Stoffer, 2026), irregular sampling has historically been handled with different techniques (see L. T. Kreutzer et al., 2023; K. Rehfeld et al., 2011): re-sampling the data onto a regular grid by interpolation before applying the standard Fourier-based estimator; binning observation pairs by their time-lag separation (“slotting”); weighting observation pairs by a smooth kernel of the time-lag separation rather than binning them; or bypassing the lag domain altogether and working with a Fourier-domain estimator adapted to irregular sampling, such as the Lomb–Scargle periodogram (J. T. VanderPlas, 2018; J. D. Scargle, 1982; N. R. Lomb, 1976).

Each of these approaches has been implemented, refined, and benchmarked by different research communities—astronomy, geoscientific context among them—largely independently of one another (Section 2). What these approaches share is a computational cost that scales at best linearly per output lag and, in the pairwise binning/weighting case, quadratically overall in the number of observations n . This is unproblematic for the short, sparse times series where these methods were originally designed for (typically $n \sim 10^2$ – 10^3), but certainly would become a practical bottleneck for the longer, denser irregular time series increasingly produced by continuous environmental monitoring, high-cadence time-domain surveys, or long instrumental records

($n \sim 10^4$ – 10^6), and for workflows that require many repeated evaluations (parameter sweeps, bootstrap resampling, batch processing of large numbers of light curves or well records).

In this paper we present `nufftcf`, an open-source Python package that computes ACF/CCF estimates by evaluating the Wiener–Khinchin theorem with the Non-Uniform Fast Fourier Transform (NUFFT) and classical-FFT for irregularly- and regularly-sampled time series respectively. `nufftcf` additionally provides direct, “real-space” implementations of the same estimators as a reference and still valuable for short time series. All estimator families share a single calling convention, so that a user can trade accuracy against speed without altering downstream analysis code.

The development of `nufftcf` yields the following contributions:

1. We demonstrate that the Gaussian- and rectangle-kernel ACF/CCF estimators—previously established on statistical grounds (K. Rehfeld et al., 2011)—can be reformulated as a pair of Non-Uniform Fast Fourier Transform (NUFFT) evaluations. This reformulation ensures numerical consistency with existing definitions while achieving an $O(n \log n)$ computational complexity.
2. We introduce an $O(n)$ analytical algorithm, based on a two-pointer scan, to compute the effective pair count for normalizing each lag bin. This replaces the naive $O(n^2)$ all-pairs approach, significantly improving efficiency.
3. We validate `nufftcf` through comparisons with `pastas` (R. A. Collenteur et al., 2019) and `pyZDCF` (I. Jankov et al., 2022b) on synthetic ACF and CCF problems.

The paper is organized as follows. Section 2 reviews prior approaches to ACF/CCF estimation for irregularly-sampled data, situates `nufftcf` relative to them, and details the methodology, from the regularly-sampled case through the four established families of irregular-sampling estimators to the NUFFT reformulation used by `nufftcf`. Section 3

compares `nufftcf` to `pastas` and `pyZDCF` on ACF and CCF use cases. Section 4 illustrates `nufftcf` on a simulated astrophysical light curve closer to observational practice. Section 5 presents timing benchmarks. Section 6 discusses limitations and guidance on estimator choice, and Section 7 concludes. The Appendix A provides additional details on the `nufftcf` algorithms and implementation choices.

2 Related Works and Methodology

This section is not intended as an exhaustive review of the correlation-function and spectral-estimation literature for irregularly-sampled data; several dedicated reviews already serve that purpose (e.g. [K. Rehfeld et al. \(2011\)](#) for kernel-weighted and interpolation-based estimators, [P. Babu & P. Stolica \(2010\)](#) for Fourier-domain spectral-analysis methods). Rather, our aim here is methodological: we introduce, in the minimal form needed to motivate and formulate `nufftcf`, the four broad families of approaches used in the literature to handle irregular sampling (Sections 2.2–2.4), before turning to the Non-Uniform Fast Fourier Transform (Section 2.5) and showing how it can be combined with kernel-weighted estimators to yield the `nufftcf` approach (Section 2.6).

2.1 Correlation estimation on a regular grid: the Wiener–Khinchin theorem

Let x_t , $t = 0, \dots, n-1$, be a zero-mean, weakly stationary process sampled at n regularly spaced times with sampling interval Δt . The (biased) sample autocovariance at lag $k\Delta t$ is

$$\hat{R}_{xx}(k) = \frac{1}{n} \sum_{t=0}^{n-k-1} x_t x_{t+k}, \quad k = 0, \dots, n-1, \quad (1)$$

and the sample ACF is¹ $\hat{\rho}_{xx}(k) = \hat{R}_{xx}(k)/\hat{R}_{xx}(0)$. The sample cross-covariance and cross-correlation be-

tween two co-sampled series x_t and y_t are defined analogously.

For weakly stationary processes, the Wiener–Khinchin theorem establishes that the autocovariance function and the power spectral density form a Fourier-transform pair. Similarly, the cross-covariance function and the cross-spectral density also form such a pair. For finite sampled data, these relationships are numerically realized through the discrete Fourier transform (DFT): the cross-spectrum is estimated as

$$\hat{S}_{xy}(f) = \hat{X}(f)\hat{Y}^*(f), \quad \hat{R}_{xy}(k) = \mathcal{F}^{-1} \left[\hat{S}_{xy}(f) \right] (k), \quad (2)$$

where $\hat{X}(f) = \mathcal{F}[x](f)$ (and similarly for $\hat{Y}(f)$) denotes the DFT of the sampled series. On a regular grid, this identity is evaluated efficiently using a pair of FFTs—a forward FFT to compute $\hat{X}(f)$ and $\hat{Y}(f)$, a pointwise complex-conjugate product to form $\hat{S}_{xy}(f)$, and an inverse FFT to recover $\hat{R}_{xy}(k)$ —yielding an $O(n \log n)$ algorithm instead of the $O(n^2)$ direct lag-domain summation of Equation 1.

This FFT-based approach provides the regular-grid baseline for correlation estimation. However, for irregularly-sampled data, where the sample times $t_1 < t_2 < \dots < t_n$ are not equally spaced, neither the sum in Equation 1 nor the FFT pair can be applied directly. The literature describes four broad strategies to address this issue, which we outline below.

2.2 Slotting and the discrete correlation function

The earliest widely used approach specific to irregular sampling is the *slotting* technique, introduced in astronomy by [R. A. Edelson & J. H. Krolik \(1988\)](#) as the discrete correlation function (DCF). Slotting bins the products of pairs of (standardized) observations by their time-lag separation into fixed-width bins and averages within each bin. While simple and model-free, the resulting correlation-function estimate is not guaranteed to be positive semi-definite and may require post-processing ([K. Rehfeld et al., 2011](#)). This

where $\hat{R}_{xx}(0)$ is the total energy of the signal. For cross-correlation, the normalization factor is $\sqrt{\hat{R}_{xx}(0)\hat{R}_{yy}(0)}$.

¹This normalization follows MATLAB convention (<https://uk.mathworks.com/help/matlab/ref/xcorr.html>),

limitation has motivated the development of kernel-weighted techniques (Section 2.3).

A refinement of the DCF, proposed by T. Alexander (1997) (ZDCF-transform), addresses the sparse, unevenly sampled light curves typical of AGN monitoring by using *equal-population binning* rather than fixed-width bins. In this approach, bin boundaries are chosen adaptively so that each bin contains approximately an equal number of pairs. Additionally, a Fisher z -transform (R. A. Fisher, 1915) is applied to stabilize the variance of the correlation estimate in each bin before computing confidence intervals. `pyZDCF` (I. Jankov et al., 2022b) is the reference Python implementation of ZDCF and remains the standard tool for CCF estimation in AGN reverberation-mapping studies. Both the DCF and ZDCF require a pairwise scan of the data, resulting in an $O(n^2)$ computational cost.

2.3 Kernel-weighted estimators

Kernel-weighted estimators replace the hard bin assignment of slotting with a smooth weighting function (kernel) of the time-lag separation. This allows pairs to contribute to a lag estimate in proportion to the proximity of their separation to that lag. In this context, P. Stoica et al. (2009) proposed a sinc kernel, while O. N. Björnstad & W. Falck (2001) introduced a Gaussian kernel. Notably, the slotting technique of R. A. Edelson & J. H. Krolik (1988) can be interpreted as a rectangle (or box) kernel.

The kernel-weighted cross-correlation estimator is defined as:

$$\hat{\rho}_{xy}(k\Delta\tau) = \frac{\sum_{i=1}^{n_x} \sum_{j=1}^{n_y} x_i y_j b_k(d)}{\sum_{i=1}^{n_x} \sum_{j=1}^{n_y} b_k(d)}, \quad (3)$$

where $d = t_j^y - t_i^x - k\Delta\tau$ represents the deviation from the target lag. Typical choices for $b_k(d)$ include the rectangle/box kernel (R. A. Edelson & J. H. Krolik, 1988):

$$b_k(d) = \begin{cases} 1 & \text{for } |d| \leq h/2, \\ 0 & \text{otherwise.} \end{cases} \quad (4)$$

and the Gaussian kernel (O. N. Björnstad & W. Falck, 2001), which is the primary choice in `pastas` and `nufftcf`:

$$b_k(d) = \frac{1}{\sqrt{2\pi}h} \exp\left(-\frac{d^2}{2h^2}\right). \quad (5)$$

Here, h is a bandwidth parameter typically set to match the mean sampling interval Δt^{xy} , see K. Rehfeld et al. (2011) for details. Another typical choice, not yet implemented in the current version of `nufftcf`, is the Hanning kernel (W. Max-Moerbeck et al., 2014).

Given the irregular sampling common in geoscientific time series, K. Rehfeld et al. (2011) conducted a systematic benchmark of kernel-based estimators (Gaussian, sinc, rectangle) against linear interpolation and the Lomb–Scargle approach, the latter being addressed in the following section. Using synthetic sinusoidal and autoregressive processes with controlled sampling irregularity, they found that the Gaussian-kernel estimator consistently exhibited the lowest, or close to the lowest, root-mean-square error (RMSE) and bias across nearly all test cases. This led the authors to recommend it over linear interpolation for irregularly sampled data.

This methodology underpins the ACF estimator implemented in `pastas` (R. A. Collenteur et al., 2019), a widely used Python package for groundwater time-series analysis. `pastas` offers Gaussian, rectangle, and regular-grid binning options for its ACF. As with slotting, kernel-weighted estimators evaluate a direct pairwise sum, with an overall cost of $O(n^2)$ for a full lag range (or $O(n)$ per individual lag).

2.4 Lomb–Scargle-based (Fourier-domain) estimation

A conceptually different approach, introduced by N. R. Lomb (1976); J. D. Scargle (1982) (see also J. D. Scargle (1989); J. T. VanderPlas (2018)), estimates the cross-correlation function by first computing the cross-spectrum via the Lomb–Scargle Fourier transform (LSFT). The LSFT is an irregular-sampling generalization of the DFT built from a least-squares sinusoid fit. The (cross-)periodogram is then formed,

and an inverse Fourier transform is applied to the result, effectively applying the Wiener–Khinchin theorem in the Fourier domain rather than directly binning or weighting observation pairs in the lag domain. More broadly, the LSFT is one instance of a wider family of Fourier-domain approaches to spectral analysis of nonuniformly-sampled data, comprehensively reviewed by [P. Babu & P. Stoica \(2010\)](#).

[K. Rehfeld et al. \(2011\)](#) found this approach competitive with kernel methods for univariate (ACF) problems but markedly worse for bivariate (CCF) problems. While the Lomb–Scargle method reuses the transform-multiply-invert idea of Section 2.1, it replaces the FFT with a nonuniform transform tailored to irregular sampling, with a computational cost of $O(n n_f)$ for n_f frequencies. In practice, this transform can itself be evaluated efficiently via a non-uniform FFT (NUFFT), following the approach of [L. H. Garrison et al. \(2024\)](#), which uses the same NUFFT library (FINUFFT) as `nufftcf`, as described in the next section.

2.5 The Non-Uniform Fast Fourier Transform

The NUFFT computes Fourier sums between nonuniformly and uniformly spaced points without forming the $O(n \cdot m)$ direct sum over all n sample points and m output frequencies (or vice versa). Two of its three standard types are relevant here (see sign convention in Section A.3):

- **type 1** (“nonuniform to uniform”): given values x_j at nonuniform points t_j , $j = 1, \dots, n$, compute

$$\hat{X}(f_k) = \sum_j x_j e^{i f_k t_j} := X_k \quad (6)$$

at m *uniform* output frequencies f_k .

- **type 2** (“uniform to nonuniform”): the adjoint operation, evaluating a uniformly-gridded Fourier series at arbitrary nonuniform output points:

$$x_j = \sum_k X_k e^{-i f_k t_j} \quad (7)$$

Modern NUFFT algorithms, such as FINUFFT ([A. H. Barnett et al., 2019](#)) (used by `nufftcf`), compute both types in $O((n+m) \log(1/\varepsilon))$ time, where ε is the requested accuracy. This is achieved by (i) *gridding*, where each nonuniform sample is spread onto a fine uniform grid using a compactly-supported interpolation kernel, (ii) taking a standard FFT of the fine grid, and (iii) *deconvolving* by dividing by the Fourier transform of the spreading kernel to correct for the smoothing introduced by the kernel. Type 2 reverses this sequence.

FINUFFT uses the “exponential of semicircle” kernel² as preading kernel, which is chosen to first minimize aliasing error for a given support width and secondly because its Fourier transform can be computed easily using a Gauss-Legendre quadrature integration. Such spreading kernel is an internal numerical-accuracy device of the NUFFT algorithm itself, distinct from the statistical weighting kernels $b_k(\cdot)$ of Equation 5. We return below to the distinction between the two kernels (spreading and statistical weighting).

2.6 From kernel-weighted estimators to NUFFT: the nufftcf approach

`nufftcf` central observation is that the kernel-weighted estimator of Equation 3 can be evaluated via Wiener–Khinchin using a NUFFT, rather than via a direct pairwise sum, whenever the weighting kernel $b_k(\cdot)$ is itself expressible as (or well approximated by) a smooth, compactly supported function of the lag—which the Gaussian and rectangle kernels used by `pastas` are. Concretely, `nufftcf`:

1. forms the (irregularly sampled) power or cross-spectrum of the input series with a **type 1 NUFFT** onto a uniform frequency grid whose resolution is set by the requested lag range and kernel bandwidth;
2. evaluates the correlation estimate at the requested lags by inverting this spectrum with a **type 2 NUFFT**, applying the Wiener–Khinchin theorem in the same spirit as the

² $\phi_\beta(x) = e^{\beta\sqrt{1-x^2}}$ for $|x| \leq 1$ (0 elsewhere).

regular-grid FFT pair of Section 2.1, but replaced by their nonuniform-capable counterparts.

Algorithms 1–3 in Appendix A give the complete pseudocode of the CCF and ACF estimators, respectively, and Appendix A.3 details the implementation choices dictated by FINUFFT conventions (domain mapping, sign convention, oversampling factor N_1 , and the interior-lag normalization used to avoid an edge-smoothing bias). The total cost of evaluating the correlation function at K lags from n (irregular) samples is $O(n \log n + K \log K)$, instead of the $O(n^2)$ (or $O(nK)$) cost of the direct sum in Equation 3. Because this route implicitly represents the signal on a finite, periodic Fourier basis, it is subject to a small spectral-leakage bias for strongly periodic input with irregular or gappy sampling (quantified in Section 5 and controllable via the number of Fourier modes). `nufftcf` therefore also provides a direct, “real-space” evaluation of the same kernel-weighted estimator, at the original $O(n^2)$ cost, both as an accuracy reference and as the recommended choice when n is modest and this bias must be avoided entirely (Section 6).

Effective pair count. Every estimate in Equation 3 is normalized by $b_k(\cdot)$ summed over all pairs—the *effective number of pairs* contributing to lag k , which both normalizes the correlation estimate and flags under-sampled lags. Computed naively this denominator is itself an $O(n^2)$ all-pairs sum. Because the sample times are sorted, `nufftcf` instead computes it in $O(n)$ with a two-pointer scan implemented as a `numba`-compiled (S. K. Lam et al., 2015) kernel per supported weighting kernel (Algorithms 4–5 in Appendix A.2), and, for the regular-grid classic-FFT estimators, recovers the same quantity for free by filtering the deterministic raw-pair-count ramp with the same discrete filter applied to the correlation numerator. This avoids reintroducing an $O(n^2)$ step alongside an otherwise $O(n \log n)$ estimator.

Unified API. All three estimator families implemented in `nufftcf`—NUFFT-based, real-space, and classic-FFT (regular grid only)—share the calling convention `fn(lags, tx, x, [ty, y], ...)`

→ `(c, b)`³, so that a user can move between them, e.g. to check a NUFFT-based result against the real-space reference on a subset of the data, without altering the rest of an analysis pipeline.

3 Comparison with `pastas` and `pyZDCF`

We now validate `nufftcf` against two widely used, independently implemented estimators — `pastas` (kernel-weighted, real-space, Section 2.3) and `pyZDCF` (equal-population slotting, Section 2.2) — on synthetic series for which the “true” correlation function is known analytically. We treat ACF and CCF separately, since `pastas` does not expose a cross-correlation estimator equivalent to its ACF⁴, so the CCF benchmark below compares only `nufftcf` against `pyZDCF`.

3.1 Test series and experimental setup

ACF benchmark series⁵. We reuse the synthetic irregularly-sampled series from the `pastas` autocorrelation-benchmark tutorial⁶, shown in Figure 1. Two underlying stationary processes are generated over a ten-year span at a one-day base resolution:

- **STS.SIN**: a pure sine wave of period $T = 365.25$ d, with analytical autocorrelation $\rho(\tau) = \cos(2\pi\tau/T)$;
- **STS.EXP**: an AR(1)-like process obtained by convolving white noise with a causal exponential

³`(c,b)` standing for “correlation estimate” and “effective pair count”.

⁴An internal cross-correlation code exists from which the ACF is computed, but it is not part of its public API. In the `pastas/pastas-plugins` repository, the cross-correlation code uses the `scipy.signal.correlate` function (P. Virtanen et al., 2020) with the FFT method, which is only valid for equidistant time steps.

⁵see `pastas_vs_nufftcf.ipynb` and `zdcf_vs_nufftcf.ipynb` notebooks.

⁶<https://pastas.readthedocs.io/stable/benchmarks/autocorrelation.html>




Figure 1: The synthetic series used in the ACF benchmark: the fully-sampled reference processes (STS_SIN, STS_EXP) and their three irregularly-sampled derivatives each (STS_SIN1-3, STS_EXP1-3), reproduced from the `pastas` autocorrelation-benchmark tutorial.

kernel $e^{-t/\alpha}$, $\alpha = 10$ d; because this kernel coincides with the process autocovariance, its analytical ACF is $\rho(\tau) = e^{-\tau/\alpha}$.

Each process is then resampled onto three irregular time grids of different sparsity, yielding six benchmark series (STS_SIN1-3, STS_EXP1-3): (1) a mix of 30-day and 14-day regular sub-sampling with dense daily coverage over the last four years; (2) the same scheme with an additional three-year data gap; and (3) re-indexing onto a set of real, highly irregular groundwater-monitoring dates spanning 1960–2015, used as is from the `pastas` tutorial. Series 3 is by far the sparsest and most irregular of the three ($n \sim 500$ –600 points over 55 years with long gaps), while series 1–2 alternate between decades of sparse sampling and a few years of dense daily coverage.

CCF benchmark series⁷. Two complementary CCF setups are used, both without access to `pastas`:

- Shifted-timestamp STS_SIN/EXP series: The same STS_SIN and STS_EXP processes used in the ACF benchmark are drawn over $n_{\text{days}} = 3650$ d, and two independent, irregular observation series $x(t_x)$ and $y(t_y)$ sample the identical under-

lying process values, with y -timestamps shifted by a known delay τ_0 relative to x -timestamps ($\sim 43\%$ of days retained in each series). Because x and y are the same signal read at shifted times, the CCF should peak exactly at $\tau = \tau_0$ with the ACF own peak value (1 for these unit-variance processes). We use $\tau_0 = 60$ d as the fiducial case (Figure 5) and additionally sweep $\tau_0 \in \{20, 60, 120\}$ d for the STS_EXP series (Figure 6) to check that the recovered peak location tracks the injected delay.

- Coupled Ornstein–Uhlenbeck (OU) processes with variable ρ . Two independent latent OU-like processes⁸ $e_1(t)$, $e_2(t)$ are generated with the same exponential-kernel construction as STS_EXP ($\alpha = 10$ d), and combined as

$$\begin{aligned} X_1(t) &= e_1(t), \\ X_2(t) &= \rho e_1(t - \tau_0) + \sqrt{1 - \rho^2} e_2(t), \end{aligned} \quad (8)$$

with $\tau_0 = 60$ d and $\rho \in \{0.1, 0.2, 0.3, 0.5, 0.7, 0.95\}$. Unlike the shifted-timestamp case, here the lag is *physical*: X_2 is a genuinely delayed, partially decorrelated copy of X_1 , contaminated by an independent noise process. The two series are sampled on independent irregular grids retaining $\sim 40\%$ and $\sim 60\%$ of days, respectively, with no artificial timestamp shift. The analytical cross-correlation function for this construction is

$$\rho_{X_1 X_2}(\tau) = \rho e^{-|\tau - \tau_0|/\alpha}, \quad (9)$$

i.e. an exponential peak of height ρ centered at $\tau = \tau_0$ (Figure 7).

⁸Concretely, each process is built by causally filtering unit-variance white noise $\eta_t \sim \mathcal{N}(0, 1)$ through a discrete exponential kernel, $e_t \propto \sum_{k \geq 0} e^{-k/\alpha} \eta_{t-k}$, truncated after a 200-day burn-in (negligible relative to $\alpha = 10$ d) and rescaled to unit variance. Since $e^{-k/\alpha}$ is geometric, this sum telescopes exactly into the AR(1) recursion $e_t = \phi e_{t-1} + \eta_t$ with $\phi = e^{-1/\alpha}$. This is the exact (not Euler–Maruyama) discretisation of the OU Stochastic Differential Equation $dX = -\theta X dt + \sigma dW$, $\theta = 1/\alpha$, at unit time step, with innovation variance $\sigma^2(1 - e^{-2\theta})/(2\theta) = 1$, i.e. $\sigma = \sqrt{2\theta/(1 - e^{-2\theta})}$, so that η_t matches the unit-variance driving noise above. The resulting process is stationary with autocorrelation $e^{-|\tau|/\alpha}$ up to the burn-in transient, which is what underlies Equation (9).

⁷see `nufftcf.ccf_demo.ipynb` notebook.

Software configuration. For `nufftcf` we evaluate both the rectangle and Gaussian kernels (Equations 4, 5) via the NUFFT-based estimator (and, for the fiducial CCF case, the direct real-space evaluation as an accuracy cross-check), using `bin_width=0.5 d` as the kernel-bandwidth control parameter for all ACF and CCF runs, and evaluating every integer lag from 1 to 365 d (ACF) or 1 to 180 d (CCF).

For `pastas` we call `pastas.stats.acf` with `bin_method` set to "rectangle" and "Gaussian" and `max_gap=30 d`, over the same lag range.

For `pyZDCF` we use `uniform_sampling = False`, `omit_zero_lags = True`, and its equal-population bin-size parameter `minpts`, the closest analogue to `nufftcf` and `pastas` bandwidth: we report `minpts = 11` (its default in practice⁹) for all CCF runs and both `minpts = 11` and `minpts = 25` for the ACF runs, with asymmetric Monte-Carlo error bars from 50 (ACF) or 100 (CCF) resamplings. Because `pyZDCF` adapts its bin boundaries to the pair count rather than to a fixed lag spacing, the number of lags it returns is set by the data and by `minpts`, i.e., not chosen by the user. It therefore always produces *far fewer* lag points than the fixed grid used by `nufftcf`, typically resulting in $\sim 100 - 500$ bins for the ACF series here, compared to the full requested grid of 365 or 180 points for `nufftcf`. This is a direct consequence of `pyZDCF` equal-population-per-bin design (Section 2.2).

3.2 Autocorrelation: `nufftcf` vs `pastas` and `pyZDCF`

Figures 2–3 compare `nufftcf` and `pastas` on the `STS_SIN` and `STS_EXP` series, respectively, showing the estimated ACF together with the residual to the analytical truth for both the Gaussian and rectangle kernels. Across all six series the two implementations are visually indistinguishable and their Root Mean Square Error (RMSE) against the analytical ACF agree to within the reported precision (e.g. RMSE = 0.03/0.05/0.04 for `STS_SIN1/2/3` and RMSE = 0.08/0.11/0.07 for `STS_EXP1/2/3`, identical for `pastas` and `nufftcf` in every case). This con-

⁹In practice, the default `minpts = 0` in `pyZDCF` is internally replaced by `ENOUGH = 11` (see `pyzdcf/pyzdcf.py`).

firms that `nufftcf`, NUFFT-based Wiener–Khinchin methodology, reproduces the direct pairwise kernel-weighted estimator of Equation 3 to the accuracy expected from its spectral-leakage bias (Section 2.6), which here is negligible compared to the sampling and estimation noise common to both methods. The residuals for both methods grow with n -dependent gaps and departures from stationarity (e.g. around the transition between sparse and dense sampling in `STS_EXP1/2` near lag ~ 40 –100 d), which is a property of the underlying data and kernel bandwidth rather than of either estimator.

Figure 4 repeats the comparison against `pyzdcf`. The `nufftcf` Gaussian and rectangle estimates, which closely follow the analytical ACF, agree with the `pyzdcf` points (within their Monte-Carlo error bars) wherever the two methods are directly comparable. The main qualitative differences are in resolution and computational cost: `nufftcf` returns a value at every requested lag for these series (e.g., 365 points) in at most a few hundred milliseconds, whereas `pyzdcf` returns only ~ 300 –500 equal-population bins (even fewer for the coarser `minpts = 25` setting) and requires several seconds per series. These timing differences are driven by `pyzdcf` $O(n^2)$ pairwise scan and Monte Carlo error resampling, compared to `nufftcf` $O(n \log n)$ NUFFT evaluation.

3.3 Cross-correlation: `nufftcf` vs `pyZDCF` (`STS_SIN/EXP` series)

Figure 5 shows the fiducial $\tau_0 = 60$ d shifted-timestamp case for the `STS_SIN` and `STS_EXP` series, respectively. The `nufftcf` Gaussian-kernel estimate is compared against its real-space reference (left panels) and its rectangle-kernel estimate (right panels). In both panels, the `pyZDCF` estimate is also displayed. The three `nufftcf` variants agree essentially exactly with one another, and all three also agree with the `pyZDCF` sparser set of binned points within its error bars. Additionally, they correctly locate the CCF peak at the injected delay $\tau = \tau_0$. Figure 6 confirms that this peak recovery is robust as τ_0 is varied over $\{20, 60, 120\}$ d, exemplified for the `STS_EXP` series: both `nufftcf` and `pyZDCF` recover peaks consistent




Figure 2: ACF of the STS_SIN1–3 series: signal (top row), `pastas` vs. `nufftcf` Gaussian-kernel ACF against the analytical truth (second row), and residuals for the Gaussian and rectangle kernels (third and fourth rows), with matching RMSE values for both estimators in every panel.




Figure 3: Same as Figure 3, for the STS_EXP1–3 series.




Figure 4: Top: ACF of the STS_SIN1–3 series: signal (top row) and ACF (bottom row), comparing `nufftcf` (Gaussian and rectangle kernels, with per-series computation time) to `pyZDCF` (`minpts`=11 and 25, with number of returned bins and computation time) and to the analytical truth. Bottom: Same for the STS_EXP11–3 series.

Table 1: Comparison of cross-correlation peaks for different values of ρ .

$\tau_0 = 60$ d	$\rho = 0.1$		$\rho = 0.2$		$\rho = 0.3$		$\rho = 0.5$		$\rho = 0.7$		$\rho = 0.95$	
	val.	loc.	val.	loc.	val.	loc.	val.	loc.	val.	loc.	val.	loc.
<code>nufftcf</code> Gaussian	0.167	167	0.255	60	0.351	60	0.539	60	0.724	60	0.962	60
<code>nufftcf</code> rectangle	0.178	167	0.258	60	0.355	60	0.546	60	0.733	60	0.973	60
realspace Gaussian	0.181	167	0.256	60	0.352	60	0.538	60	0.722	60	0.957	60
<code>pyZDCF</code>	0.072	178.7	0.082	80.9	0.094	80.9	0.117	80.9	0.138	80.9	0.154	80.9

with the true delay (marked by the vertical dashed line), with the `nufftcf` continuous curve making the peak location and shape considerably easier to read off than the `pyZDCF` handful of bins per panel.

3.4 Cross-correlation: `nufftcf` vs `pyZDCF` (coupled OU-processes)

Figure 7 presents the physically coupled Ornstein-Uhlenbeck process test of Equation 8, for which `nufftcf` (Gaussian and rectangle kernels) and `pyZDCF` estimates are displayed as well as the analytical CCF (Equation 9) as a dashed black curve.

Table 1 compares the estimated CCF peak (amplitude and location) across the six tested values of ρ , for the three estimators. The theoretical expectation is a peak value of ρ and a location of $\tau_0 = 60$ d, respectively (Eq. 9). We discuss the two peak characteristics recovery.

3.4.1 Peak location

For $\rho \geq 0.2$, both `nufftcf` kernel variants and the real-space Gaussian estimator correctly recover $\tau_0 = 60$ d, whereas `pyZDCF` systematically places the peak near 80.9 d, i.e. a constant offset of about +21 d.




Figure 5: Top: CCF of the shifted-timestamp STS.SIN test ($\tau_0 = 60$ d): `nufftcf` Gaussian vs. its own real-space reference (left) and vs. its rectangle kernel (right), both compared to `pyZDCF`. Bottom: Same legend for the shifted-timestamp STS.EXP test.




Figure 6: CCF peak recovery for the shifted-timestamp EXP network test as the injected delay τ_0 varies over $\{20, 60, 120\}$ d ($n \sim 1500$ points per network): `nufftcf` Gaussian vs. `pyZDCF`, with the true delay marked by the vertical dashed line.




Figure 7: CCF of the coupled Ornstein-Uhlenbeck processes of Equation 8 ($\tau_0 = 60$ d, $\alpha = 10$ d), for coupling strengths $\rho = 0.1, 0.4, 0.7, 0.95$: **nufftcf** (Gaussian and rectangle kernels) and **pyZDCF** against the analytical CCF of Equation 9.

This offset is an intrinsic consequence of the equal-population slotting algorithm: the reported lag is the mean lag of the pairs contained in the highest-correlation bin, not a fixed grid point, and bin edges are set purely by pair count rather than by physical lag. A test done by retaining 50% of both series (Equation 8) shows that **nufftcf** recovers the peak location, while **pyZDCF** increases the deviation by +78.1. **pyZDCF** equal-population slotting can degrade in an essentially uncontrolled way once the local pair density near τ_0 (set jointly by **minpts**, the sampling retention fractions, and their overlap structure) becomes insufficient to support a bin localized to the true peak neighborhood. We do not undertake further a detailed study of **pyZDCF** behavior as we focus ourselves on **nufftcf** results.

The $\rho = 0.1$ case is atypical for all estimators: the peak is found at lag = 167 d (**nufftcf** and real-space Gaussian alike) or 178.7 d (**pyZDCF**), far from $\tau_0 = 60$ d. At this weak coupling level, the signal-to-noise ratio of the theoretical CCF is too low for the physical peak to reliably dominate the estimated

cross-correlation function, so the global maximum is instead captured by a noise fluctuation rather than by the expected physical peak. As a further check of this interpretation, we repeated the coupled OU-process experiment with the total baseline extended by a factor of ten (from 3650 to 36500 d, i.e. ten times more sampled points at comparable cadence), keeping all other parameters unchanged. With this larger sample, **nufftcf** (both kernels) and the real-space Gaussian estimator recover lag = $\tau_0 = 60$ d even at $\rho = 0.1$. The interpretation is as the number of pairs contributing to each lag bin grows, the per-lag noise floor drops, the physical peak emerges above it, and the global maximum of the estimated CCF converges onto the true lag τ_0 .

3.4.2 Peak amplitude

A systematic, ρ -dependent bias is observed: **nufftcf** (both kernels) and the real-space Gaussian estimator all *overestimate* the peak amplitude compared to the true value, with an absolute bias that decreases with increasing ρ : e.g., +0.055/+0.058/+0.056 at

$\rho = 0.2$ versus $+0.012/+0.023/+0.007$ at $\rho = 0.95$, for `nufftcf` Gaussian and rectangle kernels, and real-space Gaussian respectively. The fact that the real-space evaluation shares essentially the same bias as the `nufftcf` NUFFT-based counterpart shows that this overestimation is inherent to the kernel-weighted estimator itself, rather than a NUFFT-specific artifact. We discuss this in more detail below. Note that the rectangle kernel consistently produces a slightly higher peak amplitude than the Gaussian kernel estimator, as the latter smooths more strongly over the peak neighborhood.

In contrast, `pyZDCF` strongly underestimates the amplitude across the entire tested range (at best estimating 0.154 instead of 0.95), with the absolute gap widening as ρ increases. This underestimation is again a direct consequence of slotting, as discussed for the peak location.

3.4.3 Systematic overestimation by the kernel estimators

The real-space Gaussian-kernel and the `nufftcf` Gaussian-kernel results agree quite well at every ρ (e.g. 0.538 vs. 0.539 at $\rho = 0.5$ and 0.957 vs. 0.962 at $\rho = 0.95$), sharing essentially the same positive bias. Since the real-space Gaussian estimator (Equation 3) and the `nufftcf` NUFFT-based estimator (Algorithm 1, step 10) numerator and denominator are weighted by the same kernel at the level of individual pairs in both implementations, the overestimation cannot be attributed to the NUFFT gridding algorithm (Section 2.5). It is instead an intrinsic property of the kernel-weighted correlation estimator itself.

We attribute this to a *selection* bias: the reported peak is a maximum, over a dense grid of lags, of a locally noisy estimate – at each lag only the pairs within a few kernel bandwidths (`bin_width` = 0.5 d) contribute, given a mean sampling interval of ~ 1.7 – 2.5 d. Scanning many such noisy per-lag estimates and reporting the global maximum systematically overestimates the true peak, an effect that is largest when ρ is comparable to the per-lag noise level and vanishes as ρ tends to 1, since the Pearson correlation is bounded by unity. This also explains the $\rho = 0.1$

case (excluded from the quantitative analysis below), where the true bump at τ_0 is not statistically distinguishable from the noise floor and the reported maximum is essentially a pure noise fluctuation, located at lag = 167 rather than 60.

To test this hypothesis quantitatively, for $\rho \geq 0.2$ and each estimator we compare on Table 2 the peak bias

$$\Delta = \text{peak val.} - \rho$$

to `std_far` the standard deviation of the CCF estimate away from the peak defined as

$$\text{std_far} = \text{std}\left(\text{CCF}(\text{lag}) \mid |\text{lag} - \tau_0| > 5\alpha = 50 \text{ d}\right)$$

The lag window is chosen so that the theoretical CCF (Equation 9) has decayed to $e^{-5} \approx 0.7\%$ of its peak value there. With lags $\in [1, 180]$ d this leaves $n_{\text{far}} = 79$ lag points. Two features support the selection bias interpretation.

First, `std_far` is essentially constant across ρ (≈ 0.038 – 0.051) for all three estimators, consistent with a signal-independent noise floor set by sampling and kernel bandwidth.

Second, $\Delta/\text{std_far}$ is of order unity at weak-to-moderate coupling (≈ 0.9 – 1.2 at $\rho = 0.2$ – 0.5) – the peak bias is comparable to the background noise, exactly as expected when selecting the maximum of a noisy curve – and decreases monotonically as ρ grows, down to ≈ 0.17 – 0.54 at $\rho = 0.95$, since `std_far` stays roughly constant while Δ itself shrinks toward the $\rho = 1$ bound. The rectangle kernel shows the largest bias-to-noise ratio at every ρ , consistent with its sharper effective bandwidth producing a noisier per-lag estimate, while real-space and `nufftcf` Gaussian remain statistically indistinguishable, confirming a kernel-choice effect rather than an implementation artifact.

This sample-size dependence can also be verified directly. Repeating the experiment with a ten-fold longer baseline (36 500 instead of 3650 d) markedly reduces $|\Delta|$ at every tested ρ , for all three estimators: e.g. $\Delta = -0.030/-0.027/-0.029$ at $\rho = 0.2$ (versus $+0.055/+0.058/+0.056$ in the fiducial run) and $-0.008/-0.001/-0.003$ at $\rho = 0.95$ (versus $+0.012/+0.023/+0.007$), confirming that the bias

Table 2: Peak bias $\Delta = \text{peak_val} - \rho$ compared to the background noise level std_far (std of the CCF estimate over $|\text{lag} - \tau_0| > 50$ d, $n_{\text{far}} = 79$), for $\rho \geq 0.2$.

Method	$\rho = 0.2$			$\rho = 0.3$			$\rho = 0.5$			$\rho = 0.7$			$\rho = 0.95$		
	Δ	std_far	ratio	Δ	std_far	ratio	Δ	std_far	ratio	Δ	std_far	ratio	Δ	std_far	ratio
nufftcf Gaussian	0.055	0.047	1.16	0.051	0.045	1.14	0.039	0.041	0.96	0.024	0.038	0.64	0.012	0.042	0.28
nufftcf rectangle	0.058	0.048	1.20	0.055	0.046	1.21	0.046	0.042	1.10	0.033	0.039	0.85	0.023	0.043	0.54
realspace Gaussian	0.056	0.051	1.10	0.052	0.048	1.07	0.038	0.044	0.88	0.022	0.040	0.55	0.007	0.044	0.17

shrinks as the per-lag noise floor drops with increasing sample size. Notably, the residual at this larger sample size turns slightly *negative* for every estimator, and is consistently less negative for the rectangle kernel than for the Gaussian kernel – the same ordering seen in the selection-bias-dominated regime – suggesting that once the selection bias is suppressed, a small residual smoothing bias from averaging the sharply peaked CCF over a finite kernel bandwidth becomes visible instead.

A natural mitigation, in both implementations (real space or NUFFT-based **nufftcf** computations), would be to widen the kernel bandwidth reducing per-lag variance at the cost of resolution or to calibrate and subtract the selection bias via a bootstrap/Monte-Carlo procedure on surrogate data, analogous to the internal Monte-Carlo simulations used by **pyZDCF** for its own uncertainty estimates. Section 4.4 revisits this same diagnosis – a single-realization deviation mistaken for a bias – on an independently-motivated, physical use case, using repeated realizations rather than a longer baseline to demonstrate it.

4 A physical use-case: Simulated Astrophysical Light Curve

The synthetic series of Section 3 were designed to isolate ACF/CCF estimator behaviour under simple, well-controlled statistical models, so as to enable a quantitative, independent cross-check against **pastas** and **pyZDCF**. To illustrate **nufftcf** on a case closer to observational practice, we simulate a single stellar photometric light curve combining

a quasi-periodic rotation signal, correlated “flicker” noise, realistic ground-based sampling gaps, and heteroscedastic measurement noise, and use **nufftcf** alone to recover the injected rotation period and noise correlation time.¹⁰

4.1 Simulated light curve

The latent flux is

$$y(t) = A(t) \cos\left(\frac{2\pi t}{P_{\text{rot}}} + \varphi\right) + n(t), \quad (10)$$

where φ is a random phase, uniform on $[0, 2\pi)$ and independent of $A(t)$; $A(t)$ is a stationary Ornstein–Uhlenbeck (OU) process of mean A_0 , variance σ_A^2 , and correlation time $\tau_A \gg P_{\text{rot}}$, modelling slow spot evolution; and $n(t)$ is an independent OU process of variance σ_n^2 and correlation time τ_n , modelling short-timescale correlated (flicker/granulation) noise. The normalized autocorrelation of $y(t)$ is then analytic,

$$\rho_{\text{true}}(\tau) = \frac{\frac{1}{2} \cos\left(\frac{2\pi\tau}{P_{\text{rot}}}\right) (\sigma_A^2 e^{-|\tau|/\tau_A} + A_0^2) + \sigma_n^2 e^{-|\tau|/\tau_n}}{\frac{1}{2}(\sigma_A^2 + A_0^2) + \sigma_n^2}, \quad (11)$$

and has the form of a sum of an exponentially damped cosine and an exponential decay, closely related to the stochastically driven damped harmonic oscillator (SHO) kernels of the celerite family (D. Foreman-Mackey et al., 2017), which have been used to model stellar rotation (e.g. R. Angus et al., 2018) and granulation/flicker noise in a similar decomposition. It is qualitatively similar in spirit, though not identical in functional form, to the product-form quasi-periodic kernels (squared-exponential times periodic) commonly used in Gaussian-process models

¹⁰See **nufftcf.astro.demo.ipynb** notebook.

Light curve: rotation + correlated noise + gaps + heteroscedastic noise




Figure 8: Simulated ground-based light curve (top; error bars show σ_i) combining stellar-rotation modulation, correlated OU noise, seasonal/weather gaps, and heteroscedastic measurement noise (bottom).

of stellar activity (e.g. V. Rajpaul et al., 2015). We use $P_{\text{rot}} = 12$ d, $A_0 = 1$, $\sigma_A = 0.35$, $\tau_A = 220$ d, $\sigma_n = 0.55$, $\tau_n = 3.5$ d. Both $A(t)$ and $n(t)$ are simulated directly at the (irregular) observation epochs using the exact OU transition kernel, so no dense-grid discretization is needed.

Ground-based sampling is emulated over four years with one nightly epoch, an ~ 8 -month seasonal observability window, and 25% of otherwise-observable nights lost to weather, yielding a fiducial realization of $n = 762$ irregular epochs (median cadence 1.0 d within a season). Each point additionally carries a heteroscedastic measurement uncertainty σ_i , growing toward the edges of the observing season (higher airmass), with $\sigma_i \in [0.07, 0.43]$ d (median 0.18) for this parameter choice; the observed flux is $x_{\text{obs}}(t_i) = y(t_i) + \varepsilon_i$, $\varepsilon_i \sim \mathcal{N}(0, \sigma_i^2)$. Figure 8 shows the resulting light curve.

4.2 ACF recovery

We compute the ACF with `nufftcf` Gaussian-kernel real-space and NUFFT estimators (`bin_width = 0.5` d) on x_{obs} , at every integer lag from 1 to 49 d. Since `nufftcf` standardizes the series before evaluating $\hat{\rho}$, and the measurement noise only contributes at lag-0 (known as “nugget” effect (e.g. A. M. Sayer, 2020) or “zero-lag” spike (e.g. B. C. Lacki, 2025)), the observed ACF is a rescaled version of Equation 11:

Recovered ACF (sparse data, heteroscedastic noise) vs true ACF




Figure 9: ACF recovered by `nufftcf` Gaussian real-space and NUFFT estimators from the sparse, heteroscedastic light curve of Figure 8, against the theoretical latent-process ACF (Equation 11).

$\rho_{\text{obs}}(\tau) = \kappa \rho_{\text{true}}(\tau)$, with dilution factor

$$\begin{aligned} \kappa &= \text{Cov}_{\text{tot}}(0) / \text{Var}(x_{\text{obs}}) \\ &= \text{Cov}_{\text{tot}}(0) / (\text{Cov}_{\text{tot}}(0) + \langle \sigma_i^2 \rangle) < 1. \end{aligned} \quad (12)$$

Figure 9 compares the two `nufftcf` estimates to ρ_{true} for the fiducial realization. Both agree closely with the theoretical curve and with each other: RMSE = 0.027 for real-space, 0.040 for NUFFT, against ρ_{true} , i.e. without correcting for the dilution factor above. The rotation-driven oscillation and the fast initial decay from the flicker noise are both clearly resolved despite the seasonal gaps and the heteroscedastic noise.

4.3 Recovering the rotation period and noise timescale

We extract two physical quantities from the recovered ACF: the rotation period P_{rot} , taken as the lag of the ACF maximum in a search window around the injected value, and the flicker-noise correlation time τ_n , obtained by fitting the short-lag (≤ 14 d) ACF to a two-component model,

$$\hat{\rho}(\tau) = \kappa \left[w \cos\left(\frac{2\pi\tau}{P_{\text{rot}}}\right) + (1 - w) e^{-\tau/\tau_n} \right], \quad (13)$$

with free parameters $(w, \tau_n, P_{\text{rot}}, \kappa)$. Note that a naive single-exponential decay does not fit this range, since $P_{\text{rot}} = 12$ d is not slowly varying compared to

the fitted lags. The explicit dilution factor κ is required in Equation 13: without it, the fit compensates for the noise-induced amplitude deficit at $\tau > 0$ by biasing τ_n low, since $\hat{\rho}$ would otherwise be pinned to 1 at $\tau = 0$ by construction.

For the fiducial realization, the ACF peak search recovers $P_{\text{rot}} = 12.00$ d exactly (injected: 12.00 d) for both estimators. The composite fit gives $\tau_n = 4.09 \pm 0.38$ d (real-space) and 4.16 ± 0.38 d (NUFFT), compared to the injected $\tau_n = 3.5$ d, and $\kappa = 0.938 \pm 0.019$ and 0.940 ± 0.020 respectively, in agreement with an independent, fit-free estimate $\kappa_{\text{est}} = 1 - \langle \sigma_i^2 \rangle / \text{Var}(x_{\text{obs}}) = 0.954$ obtained directly from the data.

4.4 Monte-Carlo over independent realizations

Taken at face value, the single-realization τ_n fit above (~ 4.1 d) is above the injected value (3.5 d) by more than its formal uncertainty, suggestive of a residual bias – reminiscent of the CCF peak-amplitude offsets identified from a single realization in Section 3.4.3. Following the same logic, we test whether this is a genuine estimator bias or a finite-sample fluctuation by repeating the full pipeline – survey sampling, heteroscedastic noise, latent signal, `nufftcf` Gaussian-NUFFT ACF, and the fit of Equation 13 – over $n_{\text{mc}} = 20$ independent realizations (Figure 10). The same model is fit to the noiseless theoretical curve, to the noise-free but irregularly-sampled latent signal, and to the fully observed (sampled+noisy) series.

Averaged over the 20 realizations, the fit recovers $\tau_n = 3.64 \pm 0.19$ d (sampling only) and 3.60 ± 0.24 d (sampling+noise) [mean $\pm$ SEM¹¹], both consistent with the injected $\tau_n = 3.5$ d, while the noiseless theoretical curve, which carries no randomness and is therefore fit only once, returns $\tau_n = 3.46$ d. The per-realization scatter is substantial (std $\simeq 0.86$ – 1.09 d), comparable to, or larger than, the ~ 0.6 d offset seen for the single fiducial realization: the latter is not statistically distinguishable from the sampling noise floor. As in the CCF peak-amplitude case of Section 3.4.3, an apparent bias identified from a sin-

¹¹SEM stands for Standard Error of the Mean.




Figure 10: Fitted τ_n over $n_{\text{mc}} = 20$ independent survey realizations, for the sampled-only latent signal and the sampled+noisy observed series, against the injected value (dashed). Medians are shown as red lines, box edges represent the first (Q1) and third (Q3) quartiles, and whiskers extend to the most extreme data point within 1.5 times the interquartile range (IQR) from the boxes.

gle, or a small number of, realization(s) is again a finite-sample/selection effect rather than an intrinsic property of the `nufftcf` estimator. This bias shrinks on average when several independent realizations are combined, even though the per-realization scatter remains unchanged, as it is determined by the survey information content rather than by the estimator itself.

This test reinforces, on an independent, physically-motivated case, the guidance of Section 3.4.3: a single-realization deviation from the truth should always be weighed against the estimator sampling variance (accessible here via repeated simulation, or in practice via bootstrap on the real data) before being interpreted as a bias.

5 Timing Benchmarks

To quantify the practical speed-up offered by `nufftcf` NUFFT, we benchmarked `nufftcf` and `pastas` on synthetic white-noise series of increas-

ing length, separately for irregularly-sampled and regularly-sampled data. **pyZDCF** was excluded from this comparison: its $O(n^2)$ real-space correlation, combined with a Monte-Carlo error estimate, makes it prohibitively slow for series of the lengths considered here – unsurprising, since **pyZDCF** was designed to prioritize error estimation over computational speed.

Benchmarks are restricted to ACF, rather than CCF, for two reasons: first, this allows a direct comparison against **pastas**, which only implements ACF; second, within **nufftcf** itself, CCF and ACF share the same computational cost, since both are driven by the same NUFFT calls and the same b -factor (pair-count normalization) computation.

For each configuration, computation time was measured over several repeats per series length, with execution order randomized across (kernel, length) combinations to avoid confounding a systematic drift (e.g. thermal throttling over a long benchmark run) with the trend in series length being measured. The reported point estimate is the per-length median (regular case) or minimum (irregular case) computation time, and error bars show the min-to-max spread across repeats. For each (kernel, algorithm) pair, timings were fit to the theoretically expected scaling law using a robust nonlinear least-squares fit (`scipy.optimize.least_squares` P. Virtanen et al. (2020), `soft_l1` loss), which down-weights the influence of occasional outlier measurements without requiring them to be manually identified and removed. Parameter uncertainties were obtained by bootstrapping over the recorded repeats (500 iterations): at each iteration, repeats at every series length are re-sampled with replacement, the point estimate is re-computed, and the model is refit; the standard deviation of the refit parameters across iterations gives the reported uncertainty.

For the sake of completeness and reproducibility, we detail the exact hardware configuration used to conduct the benchmarks: this is a MacBook Pro (2020) equipped with a 2 GHz Intel Core i5 processor (quad-core), 512 KB of L2 cache per core, and 6 MB of shared L3 cache, along with 16 GB of LPDDR4X RAM (3733 MHz, dual-channel, 8 GB per channel), running macOS 26.5.2 (build 25F84).

5.1 Irregularly-sampled data

Irregularly-sampled series are the primary use case **nufftcf** was designed for, and the NUFFT-based estimator was accordingly compared against **pastas** for both the Gaussian and rectangle smoothing kernels. Table 3 summarizes the fitted parameters, and Figure 11 shows the results. As expected from the underlying algorithms, **pastas** slotting technique scales quadratically in the number of points, $t \sim an^2 + \text{overhead}$, while **nufftcf** scales as $t \sim an \ln(n) + \text{overhead}$, consistent with its NUFFT-based evaluation of the Wiener–Khinchin relation.

The quadratic-vs- $n \ln n$ scaling translates into a rapidly growing speed advantage for **nufftcf** as series length increases, and confirms that the $O(n \log n)$ NUFFT-based approach delivers its expected asymptotic benefit in practice, not just in principle – with the small residual overhead (a few milliseconds) making **nufftcf** advantageous already for moderately sized series, well before reaching the asymptotic regime.

5.2 Regularly-sampled data

When the data lie on a uniform grid, there is no need to pay for the generality of a NUFFT: **nufftcf** also provides a dedicated classic-FFT path (`fft.acf.py`) that reproduces the exact same Gaussian/rectangle kernel definitions as the NUFFT and real-space estimators, but via a plain FFT correlation with no numba/finufft dependency in the hot path. To quantify what this dedicated fast path buys over simply reusing the general-purpose NUFFT estimator on regular data, we benchmarked **pastas** against both the FFT and NUFFT paths of **nufftcf** for the Gaussian and rectangle kernels, and additionally against **pastas** own "regular" bin method (a fixed-size windowed `numpy.corrcoef`¹² per lag, with no smoothing kernel), for which the comparison is limited to the dedicated `fft.acf` no-kernel estimator, since NUFFT has no counterpart for this case. Note that **pastas** "regular" bin method scales empirically as $O(n)$,

¹²<https://numpy.org/doc/stable/reference/generated/numpy.corrcoef.html>

Table 3: Fit results for irregularly- and regularly-sampled data. **pastas** Gaussian/rectangle bin methods and **nufftcf** nufft path follow an^2 and $an \ln(n)$ (resp.) + overhead; **pastas** "regular" bin method follows an + overhead; **nufftcf** fft path follows $an \ln(n)$ + overhead.

Sampling	Kernel	Algorithm	a	Overhead [ms]	R^2
Irregular	Gaussian	pastas	$(4.744 \pm 0.002) \times 10^{-8}$	17.32 ± 1.04	1.00
	Gaussian	nufftcf	$(1.370 \pm 0.033) \times 10^{-6}$	2.17 ± 0.59	1.00
	rectangle	pastas	$(5.448 \pm 0.004) \times 10^{-8}$	3.08 ± 0.88	1.00
	rectangle	nufftcf	$(4.712 \pm 0.230) \times 10^{-7}$	1.71 ± 0.30	0.98
Regular	Gaussian	pastas	$(4.738 \pm 0.026) \times 10^{-8}$	12.89 ± 6.31	1.00
	Gaussian	nufftcf (fft)	$(2.098 \pm 0.056) \times 10^{-8}$	1.01 ± 0.12	0.99
	Gaussian	nufftcf (nufft)	$(1.351 \pm 0.002) \times 10^{-6}$	28.43 ± 3.29	1.00
	rectangle	pastas	$(4.818 \pm 0.038) \times 10^{-8}$	0.24 ± 0.04	1.00
	rectangle	nufftcf (fft)	$(1.314 \pm 0.083) \times 10^{-8}$	0.61 ± 0.30	0.94
	rectangle	nufftcf (nufft)	$(4.267 \pm 0.073) \times 10^{-7}$	10.05 ± 2.70	0.99
	regular	pastas	$(1.461 \pm 0.016) \times 10^{-5}$	17.05 ± 0.57	1.00
	regular	nufftcf (fft)	$(1.017 \pm 0.068) \times 10^{-8}$	0.00 ± 0.06	0.99




Figure 11: ACF benchmark, **nufftcf** versus **pastas**, on irregularly-sampled data (Gaussian and rectangle kernels). Robust fit (soft_l1); error bars show the min-to-max spread across repeats at each point.

ACF benchmark (regular data): Pastas vs nufftcf (fft / nufft)




Figure 12: ACF benchmark on regularly-sampled data, `pastas` versus `nufftcf` (dedicated FFT path and general-purpose NUFFT path), Gaussian kernel.

unlike its $O(n^2)$ "Gaussian"/"rectangle" bin methods, and is fit to $t \sim an + \text{overhead}$ accordingly.

Fit results are given in Table 3, and the Gaussian kernel comparison is shown in Figure 12. The dedicated FFT path is markedly faster than NUFFT on regular data (e.g. for the Gaussian kernel, its leading coefficient a is $\sim 65\times$ smaller and its overhead $\sim 30\times$ lower), and, like the NUFFT path on irregular data, retains an $O(n \log n)$ advantage over `pastas` $O(n^2)$ Gaussian/rectangle bin methods – including over `pastas` own $O(n)$ "regular" bin method, whose much larger leading coefficient outweighs its more favorable scaling.

6 Discussion on nufftcf

6.1 Choosing an estimator: NUFFT, real-space, or classic-FFT

The three estimator families provided by `nufftcf` share a single calling convention (Section 2.6) so the choice can be made per use case. For *regularly-sampled* data, the dedicated classic-FFT (Section 5) should always be preferred: it reproduces the same Gaussian/rectangle kernel definitions at the low-

est cost, with no NUFFT-specific approximation. For *irregularly-sampled* data—the intended use case for `nufftcf`—the NUFFT estimator is the default choice, since it retains the $O(n \log n)$ scaling advantage over $O(n^2)$ slotting (Table 3) already for moderately sized series, well before the asymptotic regime is reached. However, the real-space estimator, at $O(n^2)$ cost, remains useful in the following situations: when n is modest enough (a few thousand points or fewer) that the quadratic cost is not limiting, as an reference against which a NUFFT-based result can be checked on a subset of the data.

6.2 Limitations

Two limitations, quantified in Section 3, are worth restating as guidance rather than as a defect specific to `nufftcf`. First, the Gaussian- and rectangle-kernel estimators, whether evaluated in real space or via NUFFT, share an intrinsic selection bias. This has led to overestimating a CCF peak amplitude by an amount of order the background noise level of the estimate (Table 2, Figure 7). This is a property of maximizing over many noisy per-lag estimates, not a NUFFT-specific artifact (Section 3.4.3), and it shrinks as the sampling density increases: a ten-fold

larger sample already reduces the bias by a factor of a few and leaves only a small residual smoothing bias from averaging the CCF over a finite kernel bandwidth h .

The same diagnosis extends beyond the CCF peak amplitude itself: in the physically-motivated light-curve case of Section 4.4, a single-realization time-delay fit that appeared offset from the injected value by more than its formal uncertainty turned out, once averaged over $n_{\text{mc}} = 20$ independent realizations, to be statistically consistent with the truth – confirming that the apparent bias was a finite-sample selection effect rather than an intrinsic property of the estimator.

The practical lesson is the same in both cases: a single-realization deviation from the truth should always be weighed against the estimator sampling variance – accessible via repeated simulation when available, or via bootstrap/Monte-Carlo resampling on the real data otherwise – before being interpreted as evidence of a systematic bias. In practice, this limitation is best addressed by increasing the number of sampled points or independent realizations where possible, by widening h , or by a bootstrap/Monte-Carlo calibration of the bias on surrogate data, in the spirit of `pyZDCF` own internal resampling.

Second, in contrast to `pyZDCF`, which uses coarser, data-adaptive equal-population bins, `nufftcf` uses a dense, user-chosen lag grid evaluated pointwise. This gives much higher lag resolution (Section 3.3) but, unlike `pyZDCF`, does not itself supply per-lag uncertainty estimates, which a user wanting formal error bars must still obtain externally (e.g., by bootstrap).

7 Summary and Conclusion

We have presented `nufftcf`, an open-source Python package that evaluates the Gaussian- and rectangle-kernel ACF/CCF estimators of Section 2.3 through the Wiener–Khinchin theorem using the Non-Uniform FFT from `FINUFFT` library, reducing their cost from $O(n^2)$ to $O(n \log n)$ while remaining numerically consistent with their established real-space definitions. An $O(n)$ two-pointer scan removes the remaining $O(n^2)$ bottleneck in the effective-pair-count

normalization, and dedicated classic-FFT functions extend the same estimator definitions to regularly-sampled data at a further reduced cost. All three estimator families share a single calling convention, letting a user move between speed and an exact accuracy reference without changing downstream code.

We validated `nufftcf` against `pastas` for ACF estimation and against `pyZDCF` for CCF estimation, on synthetic series with known analytical correlation structure, finding agreement with both reference implementations within their respective statistical precision (Section 3), and measured the expected $O(n \log n)$ versus $O(n^2)$ scaling advantage over `pastas` in direct timing benchmarks (Section 5). `nufftcf` is therefore positioned as a drop-in, order-of-magnitude faster alternative to existing kernel-weighted ACF/CCF estimators for the long, dense irregularly-sampled time series increasingly produced by continuous monitoring and high-cadence surveys, without sacrificing the statistical grounding of the underlying estimator.

Beyond this controlled, cross-implementation validation, we also exercised `nufftcf` on a physically-motivated simulated use case (Section 4): a stellar photometric light curve combining a quasi-periodic rotation signal, correlated “flicker” noise, realistic ground-based sampling gaps, and heteroscedastic measurement noise. `nufftcf`, used alone, recovers both the injected rotation period and the noise correlation time from the simulated observed series. Repeating this pipeline over several independent realizations further showed that an apparent bias suggested by a single realization was in fact a finite-sample selection effect rather than an intrinsic property of the estimator (Section 4.4), reinforcing on an independent, more realistic case the same conclusion reached from the synthetic peak-amplitude tests of Section 3.4.3.

Future work may include extending `nufftcf` uncertainty quantification beyond the deterministic effective pair count b following `pyZDCF` own approach, exploring additional weighting kernels within the same NUFFT reformulation. An other kind of extension may be guided by the GPU version of `FINUFFT` (Y.-h. Shih et al., 2021) for the very largest ($n \gtrsim 10^6$)

time series that may emerge from continuous environmental and astronomical monitoring.

Finally, while the quantitative validation and benchmarks in this article rely on synthetic series with known, controlled statistical properties—complemented by a single physically motivated simulated light curve—the ultimate test of `nufftcf` practical utility lies in its application to both more realistic time series simulations, such as those developed by D. Emmanoulopoulos et al. (2013), and real-world data. Tests by the AGN reverberation-mapping community currently served by `pyZDCF`, and by hydrogeologists relying on `pastas` for groundwater time series analysis, would expose `nufftcf` to observational irregularities — gaps, heteroscedastic noise, and non-stationary sampling — that synthetic series may not fully capture, and would constitute a valuable complement to the synthetic validation performed in this work.

8 Reproducible Research

The latest version of the `nufftcf` package is available at <https://github.com/jecampagne/nufftcf> under the MIT license, together with the notebooks and benchmark scripts used to produce the comparisons and timings in this article. The code is also available as a PyPI package at <https://pypi.org/project/nufftcf/> to ease its installation. The documentation is available at <https://jecampagne.github.io/nufftcf/>. For Continuous Integration the test suite is run via GitHub Actions on Linux, macOS (Apple Silicon/arm64 and Intel/x86_64), and Windows, across Python versions 3.11–3.14.

We have used `pastas` 1.14.0, while concerning `pyZDCF` we have made a fork of the <https://github.com/LSST-sersag/pyzdcf> repository to allow the notebooks to be run on the Google Colab platform¹³.

¹³<https://colab.research.google.com/>

Large Language Model Use Disclosure

For the development of `nufftcf`, we used the free version of Claude Desktop by Anthropic (Sonnet 5¹⁴, Moyen) to elaborate the Python project and set up the workflow for GitHub continuous integration and systematic tests. We also used it to assist in writing the README and documenting the notebooks and Python functions as well as their transcriptions into LaTeX algorithms.

We also used Le Chat CNRS Enterprise, an AI-powered language model developed by Mistral AI¹⁵ and tailored to the research needs of the National Centre for Scientific Research (CNRS), to refine the clarity and fluency of the English text.

All AI-generated suggestions were carefully reviewed and validated by the author.

Acknowledgments

I would like to thank Julian Hamo for fruitful discussions, as well as the developers of `FINUFFT`, the library upon which `nufftcf` is based, `pastas`, and `pyZDCF`, against which `nufftcf` is validated.

References

- Alexander, T. 1997, in *Astrophysics and Space Science Library*, Vol. 218, *Astronomical Time Series*, ed. D. Maoz, A. Sternberg, & E. M. Leibowitz, 163, doi: [10.1007/978-94-015-8941-3_14](https://doi.org/10.1007/978-94-015-8941-3_14)
- Angus, R., Morton, T., Aigrain, S., Foreman-Mackey, D., & Rajpaul, V. 2018, Inferring probabilistic stellar rotation periods using Gaussian processes, *Monthly Notices of the Royal Astronomical Society*, 474, 2094, doi: [10.1093/mnras/stx2109](https://doi.org/10.1093/mnras/stx2109)
- Babu, P., & Stoica, P. 2010, Spectral analysis of nonuniformly sampled data – a review, *Digital Signal Processing*, 20, 359, doi: <https://doi.org/10.1016/j.dsp.2009.06.019>

¹⁴<https://www.anthropic.com/news/claude-sonnet-5>.

¹⁵<https://mistral.ai>

- Barnett, A. H., Magland, J. F., & af Klinteberg, L. 2019, A parallel non-uniform fast Fourier transform library based on an “exponential of semicircle” kernel, *SIAM Journal on Scientific Computing*, 41, C479, doi: [10.1137/18M120885X](https://doi.org/10.1137/18M120885X)
- BjØrnstad, O. N., & Falck, W. 2001, Nonparametric spatial covariance functions: Estimation and testing, *Environmental and Ecological Statistics*, 8, 53, doi: [10.1023/A:1009601932481](https://doi.org/10.1023/A:1009601932481)
- Collenteur, R. A., Bakker, M., Caljé, R., Klop, S. A., & Schaars, F. 2019, Pastas: Open Source Software for the Analysis of Groundwater Time Series, *Groundwater*, 57, 877, doi: [10.1111/gwat.12925](https://doi.org/10.1111/gwat.12925)
- Edelson, R. A., & Krolik, J. H. 1988, The Discrete Correlation Function: A New Method for Analyzing Unevenly Sampled Variability Data, *ApJ*, 333, 646, doi: [10.1086/166773](https://doi.org/10.1086/166773)
- Emmanoulopoulos, D., McHardy, I. M., & Papadakis, I. E. 2013, Generating artificial light curves: revisited and updated, *MNRAS*, 433, 907, doi: [10.1093/mnras/stt764](https://doi.org/10.1093/mnras/stt764)
- Fisher, R. A. 1915, Frequency distribution of the values of the correlation coefficient in samples from an indefinitely large population, *Biometrika*, 10, 507, doi: [10.2307/2331838](https://doi.org/10.2307/2331838)
- Foreman-Mackey, D., Agol, E., Ambikasaran, S., & Angus, R. 2017, Fast and Scalable Gaussian Process Modeling with Applications to Astronomical Time Series, *The Astronomical Journal*, 154, 220, doi: [10.3847/1538-3881/aa9332](https://doi.org/10.3847/1538-3881/aa9332)
- Garrison, L. H., Foreman-Mackey, D., Shih, Y.-h., & Barnett, A. 2024, NIFTY-LS: Fast and Accurate Lomb-Scargle Periodograms Using a Non-uniform FFT, *Research Notes of the American Astronomical Society*, 8, 250, doi: [10.3847/2515-5172/ad82cd](https://doi.org/10.3847/2515-5172/ad82cd)
- Jankov, I., Kovačević, A. B., Ilić, D., et al. 2022a, Photoreverberation mapping of quasars in the context of Legacy Survey of Space and Time observing strategies, *Astronomische Nachrichten*, 343, e210090, doi: [10.1002/asna.20210090](https://doi.org/10.1002/asna.20210090)
- Jankov, I., Kovačević, A. B., Ilić, D., Sánchez-Sáez, P., & Nikutta, R. 2022b, pyZDCF: Initial Release, doi: [10.5281/zenodo.7253034](https://doi.org/10.5281/zenodo.7253034)
- Kovačević, A., Popović, L. Č., Shapovalova, A. I., et al. 2014, Time series analysis of active galactic nuclei: The case of Arp 102B, 3C 390.3, NGC 5548 and NGC 4051, *Advances in Space Research*, 54, 1414, doi: [10.1016/j.asr.2014.06.025](https://doi.org/10.1016/j.asr.2014.06.025)
- Kreutzler, L. T., Gillen, E., Briegal, J. T., & Queloz, D. 2023, S-ACF: a selective estimator for the autocorrelation function of irregularly sampled time series, *MNRAS*, 522, 5049, doi: [10.1093/mnras/stad1223](https://doi.org/10.1093/mnras/stad1223)
- Lacki, B. C. 2025, Flickers, Bursts, and Dips: Detecting Rapid Variability with the g(2) Autocorrelation Function, *Astrophys. J.*, 987, 68, doi: [10.3847/1538-4357/add151](https://doi.org/10.3847/1538-4357/add151)
- Lam, S. K., Pitrou, A., & Seibert, S. 2015, in *Proceedings of the Second Workshop on the LLVM Compiler Infrastructure in HPC, LLVM '15* (New York, NY, USA: Association for Computing Machinery), doi: [10.1145/2833157.2833162](https://doi.org/10.1145/2833157.2833162)
- Lomb, N. R. 1976, Least-Squares Frequency Analysis of Unequally Spaced Data, *Ap&SS*, 39, 447, doi: [10.1007/BF00648343](https://doi.org/10.1007/BF00648343)
- Max-Moerbeck, W., Richards, J. L., Hovatta, T., et al. 2014, A method for the estimation of the significance of cross-correlations in unevenly sampled red-noise time series, *MNRAS*, 445, 437, doi: [10.1093/mnras/stu1707](https://doi.org/10.1093/mnras/stu1707)
- Rajpaul, V., Aigrain, S., Osborne, M. A., Reece, S., & Roberts, S. 2015, A Gaussian process framework for modelling stellar activity signals in radial velocity data, *Monthly Notices of the Royal Astronomical Society*, 452, 2269, doi: [10.1093/mnras/stv1428](https://doi.org/10.1093/mnras/stv1428)
- Rehfeld, K., Marwan, N., Heitzig, J., & Kurths, J. 2011, Comparison of correlation analysis techniques for irregularly sampled time series, *Nonlinear Processes in Geophysics*, 18, 389, doi: [10.5194/npg-18-389-2011](https://doi.org/10.5194/npg-18-389-2011)

- Sánchez-Sáez, P., Lira, H., Martí, L., et al. 2021, Searching for Changing-state AGNs in Massive Data Sets. I. Applying Deep Learning and Anomaly-detection Techniques to Find AGNs with Anomalous Variability Behaviors, *AJ*, 162, 206, doi: [10.3847/1538-3881/ac1426](https://doi.org/10.3847/1538-3881/ac1426)
- Sayer, A. M. 2020, How Long Is Too Long? Variogram Analysis of AERONET Data to Aid Aerosol Validation and Intercomparison Studies, *Earth and Space Science*, 7, e2020EA001290, doi: <https://doi.org/10.1029/2020EA001290>
- Scargle, J. D. 1982, Studies in astronomical time series analysis. II. Statistical aspects of spectral analysis of unevenly spaced data., *ApJ*, 263, 835, doi: [10.1086/160554](https://doi.org/10.1086/160554)
- Scargle, J. D. 1989, Studies in Astronomical Time Series Analysis. III. Fourier Transforms, Autocorrelation Functions, and Cross-Correlation Functions of Unevenly Spaced Data, *ApJ*, 343, 874, doi: [10.1086/167757](https://doi.org/10.1086/167757)
- Shapovalova, A. I., Popović, L. Č., Afanasiev, V. L., et al. 2019, Long-term optical spectral monitoring of a changing-look active galactic nucleus NGC 3516 - I. Continuum and broad-line flux variability, *MNRAS*, 485, 4790, doi: [10.1093/mnras/stz692](https://doi.org/10.1093/mnras/stz692)
- Shih, Y.-h., Wright, G., Andén, J., Blaschke, J., & Barnett, A. H. 2021, cuFINUFFT: a load-balanced GPU library for general-purpose nonuniform FFTs, *arXiv e-prints*, arXiv:2102.08463, doi: [10.48550/arXiv.2102.08463](https://doi.org/10.48550/arXiv.2102.08463)
- Stoffer, D. 2026, *astsa: Applied Statistical Time Series Analysis*. <https://CRAN.R-project.org/package=astsa>
- Stoica, P., Li, J., & He, H. 2009, Spectral analysis of nonuniformly sampled data: a new approach versus the periodogram, *Trans. Sig. Proc.*, 57, 843–858, doi: [10.1109/TSP.2008.2008973](https://doi.org/10.1109/TSP.2008.2008973)
- VanderPlas, J. T. 2018, Understanding the Lomb-Scargle Periodogram, *ApJS*, 236, 16, doi: [10.3847/1538-4365/aab766](https://doi.org/10.3847/1538-4365/aab766)
- Virtanen, P., Gommers, R., Oliphant, T. E., et al. 2020, SciPy 1.0: Fundamental Algorithms for Scientific Computing in Python, *Nature Methods*, 17, 261, doi: [10.1038/s41592-019-0686-2](https://doi.org/10.1038/s41592-019-0686-2)
- Welsh, W. 2024, in *AAS/High Energy Astrophysics Division*, Vol. 21, AAS High Energy Astrophysics Division Meeting #21, 102.18

A NUFFT-based ACF/CCF algorithms

This appendix provides the complete pseudocode for the NUFFT-based estimators implemented using the FINUFFT library, as outlined in Section 2.6 (`nufft_ccf.py` and `nufft_acf.py`). These estimators evaluate the Wiener–Khinchin identity of Equation 2 via a type 1 NUFFT (from nonuniform samples to uniform frequency grid) followed by a type 2 NUFFT (from uniform frequency grid to requested lags). Both weighting kernels supported by `nufftcf` (Gaussian and rectangle; see Equations 4–5) share the same algorithmic structure. They differ only in the smoothing operator $\mathcal{S}_h$ (Gaussian filter of width h , or box filter of size $\text{round}(2h)$) applied to the raw NUFFT output, and in the corresponding pair-count function $b(\cdot)$ from Appendix A.2. Thus, each estimator is presented as a single, kernel-parameterized algorithm to avoid redundancy.

A.1 Cross and Autocorrelation functions

Algorithm 1 describes the general NUFFT-based CCF estimator. It evaluates the Wiener–Khinchin identity by applying a type 1 NUFFT to each series, computes the cross-spectrum $\hat{X}^*(f)\hat{Y}(f)$, evaluates it at the requested lags using a type 2 NUFFT, smooths the result with the chosen kernel $\mathcal{S}_h$, and normalizes it to a Pearson correlation in $[-1, 1]$. This normalization uses the matching effective pair count (Algorithm 4) and the geometric mean of the two series ACF values at lag 0, $\sqrt{\text{ACF}_x(0) \cdot \text{ACF}_y(0)}$, each obtained from Algorithm 2.

Algorithm 1: CCF via NUFFT + smoothing (Gaussian or rectangle)

Require : times $t = (t_i^x)_{i=1}^{n_x}$, values $x = (x_i)$ (series 1, t ascending); times $s = (t_j^y)_{j=1}^{n_y}$, values $y = (y_j)$ (series 2, s ascending); requested lags $\{\tau_1, \dots, \tau_K\}$; kernel filter S_h with bandwidth h (σ for the Gaussian kernel, half-width for the rectangle kernel); NUFFT tolerance ε ; frequency-grid size N_1 (default $32 \max(n_x, n_y)$)

Ensure : $c = (c_1, \dots, c_K)$ (CCF estimate, normalized to $[-1, 1]$), $b = (b_1, \dots, b_K)$ (effective pair count)

```

1  $x_{\text{std}} \leftarrow (x - \bar{x})/\text{std}(x)$ ,  $y_{\text{std}} \leftarrow (y - \bar{y})/\text{std}(y)$ ; ▷ standardize both series
2  $t_{\min} \leftarrow \min(\min t, \min s)$ ,  $\text{span} \leftarrow \max(\max t, \max s) - t_{\min}$ ; ▷ common (union) time span
3  $\tilde{t} \leftarrow \frac{t - t_{\min}}{\text{span}} 2\pi$ ,  $\tilde{s} \leftarrow \frac{s - t_{\min}}{\text{span}} 2\pi$ ,  $\tilde{\tau} \leftarrow \frac{(0, \tau_1, \dots, \tau_K)}{\text{span}} 2\pi$ ; ▷ lag 0 prepended for the normalization
   step below
4  $\hat{F}_1 \leftarrow \text{NUFFT1D1}(\tilde{t}, x_{\text{std}}, N_1, \varepsilon)$ ; ▷ type 1: time  $\rightarrow$  frequency
5  $\hat{F}_2 \leftarrow \text{NUFFT1D1}(\tilde{s}, y_{\text{std}}, N_1, \varepsilon)$ ;
6  $M \leftarrow \hat{F}_1 \odot \hat{F}_2$ ; ▷ cross-spectrum  $\hat{X}^*(f)\hat{Y}(f)$ , not Hermitian in general
7  $c_{\text{raw}} \leftarrow \Re[\text{NUFFT1D2}(\tilde{\tau}, M, \varepsilon)]$ ; ▷ type 2: frequency  $\rightarrow$  lags
8  $c_{\text{sm}} \leftarrow S_h(c_{\text{raw}})$ ; ▷ Gaussian or rectangle filtering
9  $b_{\text{cross}} \leftarrow \text{PAIRCOUNT}(t, s, (0, \tau_1, \dots, \tau_K), h)$ ; ▷ Alg. 4, matching kernel
10  $c_{\text{norm}} \leftarrow c_{\text{sm}} \oslash b_{\text{cross}}$ ; ▷ elementwise division
11  $\text{scale}_x \leftarrow \text{ACFSCALEATLAG0}(t, x_{\text{std}}, t_{\min}, \text{span}, N_1, \varepsilon, h)$ ; ▷ Alg. 2
12  $\text{scale}_y \leftarrow \text{ACFSCALEATLAG0}(s, y_{\text{std}}, t_{\min}, \text{span}, N_1, \varepsilon, h)$ ;
13  $\text{scale} \leftarrow \sqrt{\text{scale}_x \cdot \text{scale}_y}$ ;
14  $c \leftarrow c_{\text{norm}}[2:] / \text{scale}$ ; ▷ drop the lag-0 entry (indexing starts at 1)
15  $b \leftarrow b_{\text{cross}}[2:]$ ;
16 return  $c, b$ 

```

The ACF (Algorithm 3) is a special case of this estimator where $s = t$ and $y = x$. This simplifies the procedure in three ways. First, since only one series is involved, the power spectrum $|\hat{X}(f)|^2$ is real and even, so only one type 1 NUFFT call is required. The $\pm i$ sign convention, which must be carefully fixed for the CCF cross-spectrum, is irrelevant here. Second, the effective pair count reduces to Algorithm 4 evaluated with $s = t$. Third, normalization is simpler: instead of computing the geometric mean of two scales $\text{ACF}_x(0) \cdot \text{ACF}_y(0)$ separately (Algorithm 2), the ACF estimate is normalized directly by the smoothed value at lag 0, obtained as the first entry of the same evaluation array. The interior-position trick of Algorithm 2 is unnecessary here because the ACF is even about $\tau = 0$.

A.2 Effective pair count algorithms

The pair-count functions $b(\cdot)$ (module `kernels.py`) are presented in Algorithms 4-5. Each algorithm computes the *effective* number of pairs contributing to each lag — the normalization denominator in Equation 3. They are implemented in `numba` using a two-pointer scan: since both time arrays are sorted in ascending order, the window $[\text{lo}, \text{hi}]$ associated with the pair-matching center advances monotonically as the outer index increases. This reduces the computational cost from $O(n_1 n_2)$ to $O(n_1 + n_2)$ per lag, and to an overall $O((n_1 + n_2) \log n_1)$ -equivalent throughput when parallelized (`prange`) across lags.

The ACF and CCF pair-counts are *not* separate algorithms but two instantiations of the same two-pointer scan, differing only in the role of the time arrays and the sign $\epsilon \in \{+1, -1\}$ in the window-center formula $c_j = t_j^y + \epsilon \tau_k$. For the ACF ($t^y = t^x = t$, $\epsilon = +1$), the window center is $c_j = t_j +$

Algorithm 2: ACF scale at lag 0 (Pearson normalization) — ACFSCALEATLAG0

Require : $t, x_{\text{std}}, t_{\min}, \text{span}$ (shared by both series), N_1, ε , bandwidth h
Ensure : scalar estimate of $\text{ACF}_x(0)$

- 1 $\tilde{t} \leftarrow (t - t_{\min}) / \text{span} \cdot 2\pi$;
- 2 $\hat{F} \leftarrow \text{NUFFT1D1}(\tilde{t}, x_{\text{std}}, N_1, \varepsilon)$, $P \leftarrow |\hat{F}|^2$;
- 3 $n_h \leftarrow \max(\lceil 4h \rceil + 2, 5)$; $\tau_{\text{sym}} \leftarrow (-n_h, \dots, -1, 0, 1, \dots, n_h)$; ▷ lag 0 placed at an *interior* array position
- 4 $\tilde{\tau}_{\text{sym}} \leftarrow \tau_{\text{sym}} / \text{span} \cdot 2\pi$;
- 5 $c_{\text{raw}} \leftarrow \Re[\text{NUFFT1D2}(\tilde{\tau}_{\text{sym}}, P, \varepsilon)]$;
- 6 $c_{\text{sm}} \leftarrow \mathcal{S}_h(c_{\text{raw}})$; ▷ same smoothing as Alg. 1, applied symmetrically
- 7 $b_{\text{sym}} \leftarrow \text{PAIRCOUNT}(t, t, \tau_{\text{sym}}, h)$; ▷ Alg. 4 with $s = t$
- 8 **return** $c_{\text{sm}}[n_h + 1] / \max(b_{\text{sym}}[n_h + 1], 10^{-16})$

Algorithm 3: ACF via NUFFT + smoothing (Gaussian or rectangle)

Require : times t , values x ; lags $\{\tau_1, \dots, \tau_K\}$; bandwidth h ; ε ; N_1 (default $32n$)
Ensure : $c = (c_1, \dots, c_K)$, $b = (b_1, \dots, b_K)$

- 1 $x_{\text{std}} \leftarrow (x - \bar{x}) / \text{std}(x)$;
- 2 $\tilde{t} \leftarrow \frac{t - \min t}{\max t - \min t} 2\pi$, $\tilde{\tau} \leftarrow \frac{(0, \tau_1, \dots, \tau_K)}{\max t - \min t} 2\pi$;
- 3 $\hat{F} \leftarrow \text{NUFFT1D1}(\tilde{t}, x_{\text{std}}, N_1, \varepsilon)$;
- 4 $M \leftarrow \hat{F} \odot \overline{\hat{F}} = |\hat{F}|^2$; ▷ real and even: the Hermitian case, no sign ambiguity
- 5 $c_{\text{raw}} \leftarrow \Re[\text{NUFFT1D2}(\tilde{\tau}, M, \varepsilon)]$;
- 6 $c_{\text{sm}} \leftarrow \mathcal{S}_h(c_{\text{raw}})$;
- 7 $b_{\text{eval}} \leftarrow \text{PAIRCOUNT}(t, t, (0, \tau_1, \dots, \tau_K), h)$; ▷ Alg. 4, $s = t$
- 8 $c_{\text{eval}} \leftarrow c_{\text{sm}} \odot b_{\text{eval}}$;
- 9 $c \leftarrow c_{\text{eval}}[2:] / c_{\text{eval}}[1]$; ▷ normalize by the lag-0 value
- 10 $b \leftarrow b_{\text{eval}}[2:]$;
- 11 **return** c, b

Algorithm 4: PAIRCOUNT — Gaussian (weighted) kernel, generic (ACF/CCF)

Require : $t^y = (t_j^y)_{j=1}^{n_y}$ sorted (outer array), $t^x = (t_i^x)_{i=1}^{n_x}$ sorted (window array), lags $\{\tau_k\}_{k=1}^K$, bandwidth h ,
 sign $\epsilon \in \{+1, -1\}$
Ensure : $b = (b_k)_{k=1}^K$

```

1 for  $k \leftarrow 1$  to  $K$  in parallel
2    $lo \leftarrow 0, hi \leftarrow 0, b_k \leftarrow 0$ ;
3   for  $j \leftarrow 1$  to  $n_y$  do
4      $\triangleright$  increasing  $j \Rightarrow$  increasing center  $\Rightarrow$  lo/hi only advance
5     center  $\leftarrow t_j^y + \epsilon \tau_k$ ;
6     while  $lo < n_x$  and  $t_{lo}^x < \text{center} - 6h$  do
7        $lo \leftarrow lo + 1$ ;
8     while  $hi < n_x$  and  $t_{hi}^x < \text{center} + 6h$  do
9        $hi \leftarrow hi + 1$ ;
10    for  $i \leftarrow lo$  to  $hi - 1$  do
11       $b_k \leftarrow b_k + \frac{1}{h\sqrt{2\pi}} \exp\left(-\frac{(t_i^x - \text{center})^2}{2h^2}\right)$ ;
12   $b_k \leftarrow \max(b_k, 10^{-16})$ ;  $\triangleright$  avoids division by zero at sparsely-populated lags
13 return  $b$  • ACF:  $t^y=t^x=t, n_y=n_x=n, \epsilon=+1 \Rightarrow \text{compute\_b\_Gaussian.}$  • CCF:  $t^y=s, t^x=t, \epsilon=-1 \Rightarrow$   

    $\text{compute\_b\_Gaussian\_cross.}$ 

```

τ_k with pair condition $t_i - t_j \approx \tau_k$, corresponding to `compute_b.Gaussian` and `compute_b.rectangle`. For the CCF ($t^y = s, t^x = t, \epsilon = -1$), the window center is $c_j = s_j - \tau_k$ with pair condition $s_j - t_i \approx \tau_k$, corresponding to `compute_b.Gaussian_cross` and `compute_b.rectangle_cross`.

The sign flip encodes the asymmetry of the CCF (lag $+\tau$ means y lags x by τ ; cf. the sign-convention note on the cross-spectrum). For the ACF, the sign choice ($t_i - t_j \approx \tau_k$ rather than $t_j - t_i \approx \tau_k$) is an arbitrary but fixed convention, harmless because x is correlated with itself.

A.3 Implementation choices dictated by FINUFFT

Algorithms 1–3 rely on the FINUFFT library for both nonuniform transforms (`nufft1d1`, type 1, and `nufft1d2`, type 2). Several implementation choices follow directly from this library constraints and conventions. We detail them here, complementing the $O(n)$ -per-lag numba kernels of Appendix A.2.

Time domain $[0, 2\pi)$. FINUFFT requires nonuniform points to lie in the standard domain $[-\pi, \pi)$ or $[0, 2\pi)$, so the user-supplied sample times t (in days, or any other unit) must first be rescaled. For the ACF, a single series is involved and its own extent is used:

$$\tilde{t} = \frac{t - \min(t)}{\max(t) - \min(t)} 2\pi.$$

For the CCF, however, *two* series t and s , potentially shifted and with different supports, must be placed on the same frequency grid for the cross-product $\overline{\hat{F}_1} \odot \hat{F}_2$ to be meaningful. Algorithm 1 therefore uses the *union* of the two time spans, $t_{\min} = \min(\min(t), \min(s))$, $\text{span} = \max(\max(t), \max(s)) - t_{\min}$, and applies the same affine map to t, s , and to the requested lags τ . Because lags have units of time, they are rescaled identically, $\tilde{\tau} = 2\pi\tau/\text{span}$, which is what allows the type 2 output to be read off directly at the desired lags without any further inverse transform.

Sign convention and the non-Hermitian cross-spectrum. With FINUFFT default sign conventions,

Algorithm 5: PAIRCOUNT — rectangle (counting) kernel, generic (ACF/CCF)

Require : $t^y = (t_j^y)_{j=1}^{n_y}$ sorted (outer array), $t^x = (t_i^x)_{i=1}^{n_x}$ sorted (window array), lags $\{\tau_k\}_{k=1}^K$, half-width h ,
 sign $\epsilon \in \{+1, -1\}$
Ensure : $b = (b_k)_{k=1}^K$

```

1 for  $k \leftarrow 1$  to  $K$  in parallel
2   lo  $\leftarrow 0$ , hi  $\leftarrow 0$ ,  $b_k \leftarrow 0$ ;
3   for  $j \leftarrow 1$  to  $n_y$  do
4     center  $\leftarrow t_j^y + \epsilon \tau_k$ ;
5     while lo  $< n_x$  and  $t_{lo}^x < \text{center} - h$  do
6       lo  $\leftarrow \text{lo} + 1$ ;
7     hi  $\leftarrow \max(\text{hi}, \text{lo})$ ;
8     while hi  $< n_x$  and  $t_{hi}^x \leq \text{center} + h$  do
9       hi  $\leftarrow \text{hi} + 1$ ;
10     $b_k \leftarrow b_k + (\text{hi} - \text{lo})$ ; ▷ direct count, no weighted sum:  $O(1)$  per  $j$ 
11   $b_k \leftarrow \max(b_k, 10^{-16})$ ;
12 return  $b$  • ACF:  $t^y=t^x=t$ ,  $\epsilon=+1 \Rightarrow \text{compute\_b\_rectangle}$ . • CCF:  $t^y=s$ ,  $t^x=t$ ,  $\epsilon=-1 \Rightarrow$   

   compute\_b\_rectangle\_cross.

```

isign=+1 for type-1 (Equation 6) and isign=-1 for type-2 (Equation 7), this corresponds to the usual opposite-sign Fourier-transform conventions. For the ACF, the power spectrum $|\hat{X}(f)|^2$ is real and even, so the sign convention has no effect: $M = \hat{F} \odot \overline{\hat{F}}$. For the CCF, the cross-spectrum $\hat{X}^*(f)\hat{Y}(f)$ is Hermitian (since x, y are real-valued), which guarantees that c_{raw} is real up to NUFFT roundoff. However, it does not generally have exchange symmetry $x \leftrightarrow y$:

$$S_{xy}(-f) = S_{yx}(f) \neq S_{xy}(f) \text{ in general,}$$

mirroring the time-domain identity $C_{yx}(\tau) = C_{xy}(-\tau)$. Consequently, with the convention $C_{xy}(\tau) = \int x(t)y(t+\tau)dt$, the choice $M = \hat{F}_1 \odot \hat{F}_2$ places the correlation peak at the physical lag $+\tau$ when signal 2 is delayed by $+\tau$ relative to signal 1, whereas the opposite choice $\hat{F}_1 \odot \overline{\hat{F}_2}$ effectively computes $S_{yx}(f)$ and therefore places the peak at $-\tau$, i.e. a time-reversed (mirrored) result. This is a matter of Fourier/CCF convention rather than numerical precision, but the convention must be applied consistently for the sign of the estimated delay to have the intended physical meaning.

Frequency-grid size N_1 . The parameter N_1 (Algorithms 1-3) sets the resolution of the type 1 transform (time $\rightarrow$ frequency), and therefore the fidelity of the reconstructed signal ahead of the type 2 evaluation. The default $N_1 = 32 \max(n_x, n_y)$ (a $32\times$ oversampling factor) was validated empirically against the exact real-space estimator (`compute_ccf.Gaussian_realspace`; see `kernels.py`): a factor of 32 is enough to bring the two estimators into close agreement for both the Gaussian and rectangle kernels. Higher oversampling yields only marginal further improvement (mainly for strongly periodic signals with the Gaussian kernel) at the cost of increased $O(N_1 \log N_1)$ time and memory. Scaling $N_1 \propto n$, rather than fixing a constant grid size, keeps the frequency grid adequate regardless of series length.

Requested tolerance ϵ . The `eps` parameter sets the accuracy requested from FINUFFT internal iterative scheme (gridding + FFT + deconvolution; Section 2.5). While FINUFFT own library default is 10^{-6} , our implementation uses a tighter default of 10^{-9} (Algorithms 1-3). Looser values such as 10^{-6} speed up

the computation but introduce NUFFT-specific numerical noise.

Lag-0 normalization and interior placement.

The smoothing step (Gaussian or box filter) applied to the raw NUFFT output (Algorithm 1 step 8, Algorithm 3 step 6) introduces an edge artifact: a point at the *end* of the evaluated array is not smoothed symmetrically, unlike an *interior* point. The CCF peak of interest (at lag τ_0 , interior to the requested lags range) *is* an interior point. Dividing this peak by an $\text{ACF}(0)$ scale computed by placing lag 0 at the *first* position of an array (hence smoothed asymmetrically) introduced a systematic 2–3% deficit in the estimated CCF peak. Algorithm 2 (step 3) corrects this by evaluating $\text{ACF}_x(0)$ on a small, *symmetric* lag array centered on 0 (indices $-n_h, \dots, 0, \dots, n_h$, with $n_h = \max(\lceil 4h \rceil + 2, 5)$, i.e. at least $4h$ of margin on each side), so that the smoothing filter acts there exactly as it does at the peak — a concrete example of a numerical-pipeline artifact (not a property of the underlying signal) that must be neutralized by how the evaluation array is constructed, rather than corrected *a posteriori*.

Prior standardization. Both x and y are standardized (zero mean, unit variance) before any call to FINUFFT. This is not dictated by FINUFFT itself, but ensures that the NUFFT internal scale (which depends on input amplitude) is consistent between the two compared series; combined with the final division by $\sqrt{\text{ACF}_x(0) \cdot \text{ACF}_y(0)}$, it guarantees $c \in [-1, 1]$ under the Pearson convention, regardless of the original physical scale of either series.